Track your Claude Code usage

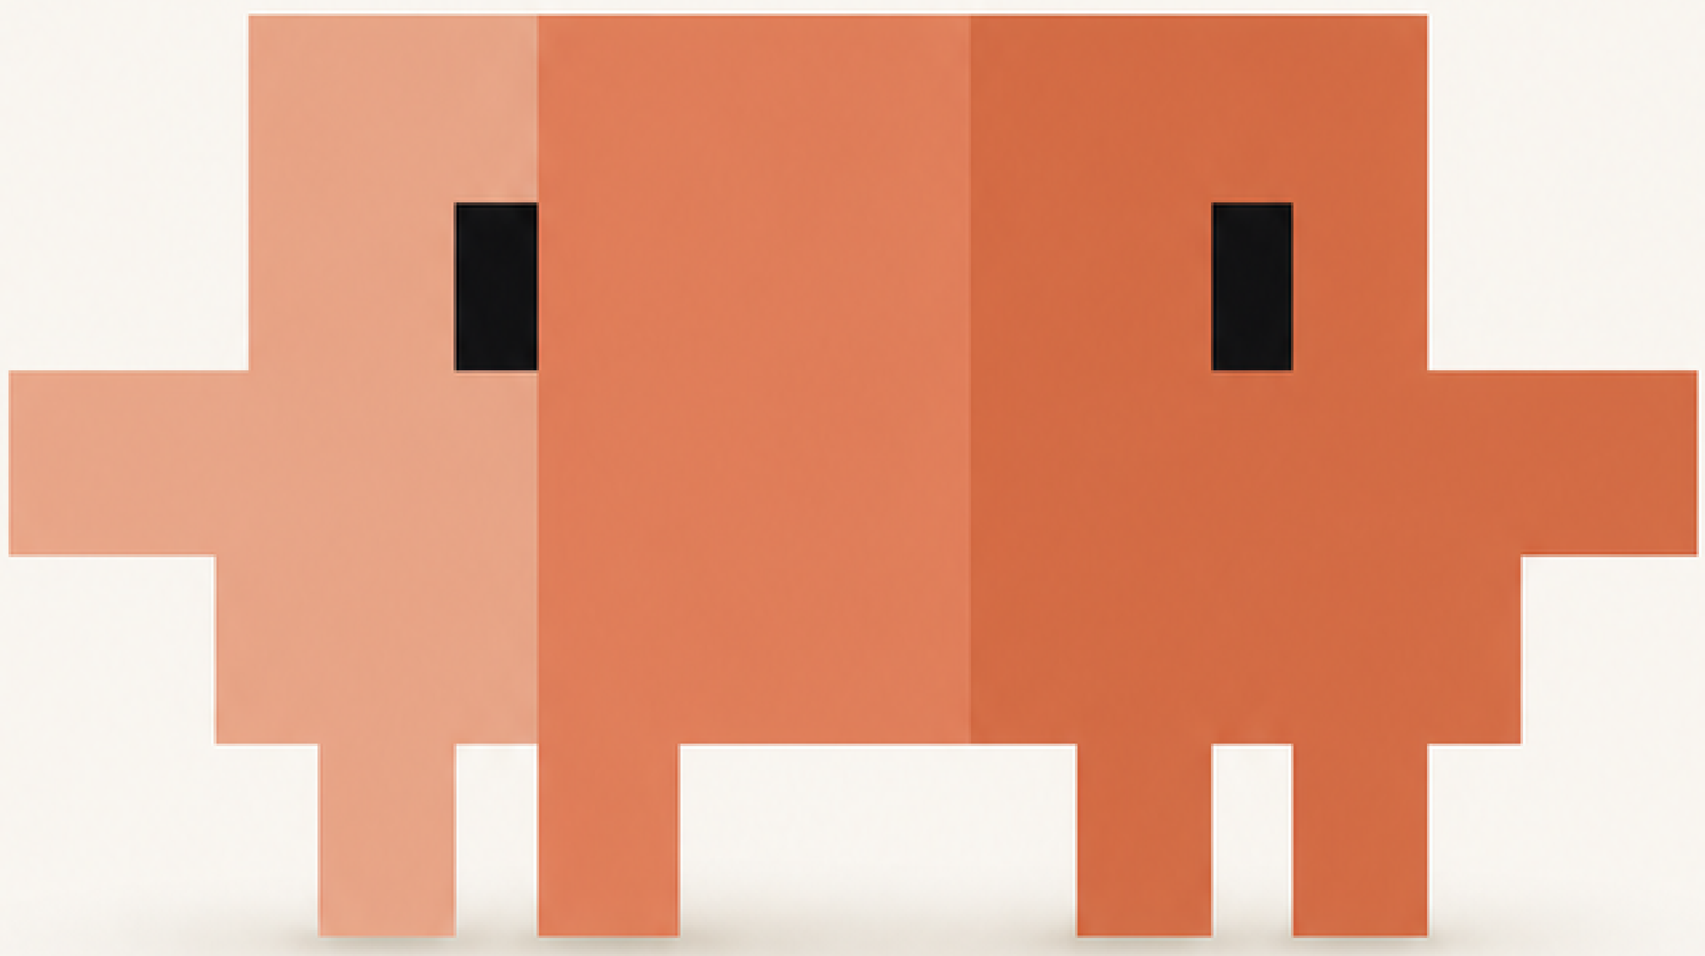

LangWatch

Cache hit tracking

Cache reads and writes are separate token classes on every span, down to the 5m and 1h TTL split on writes.

Theoretical vs billed costs

Exact cost numbers, straight from the API response. See how much you would be spending if you weren't on an individual or team plan.

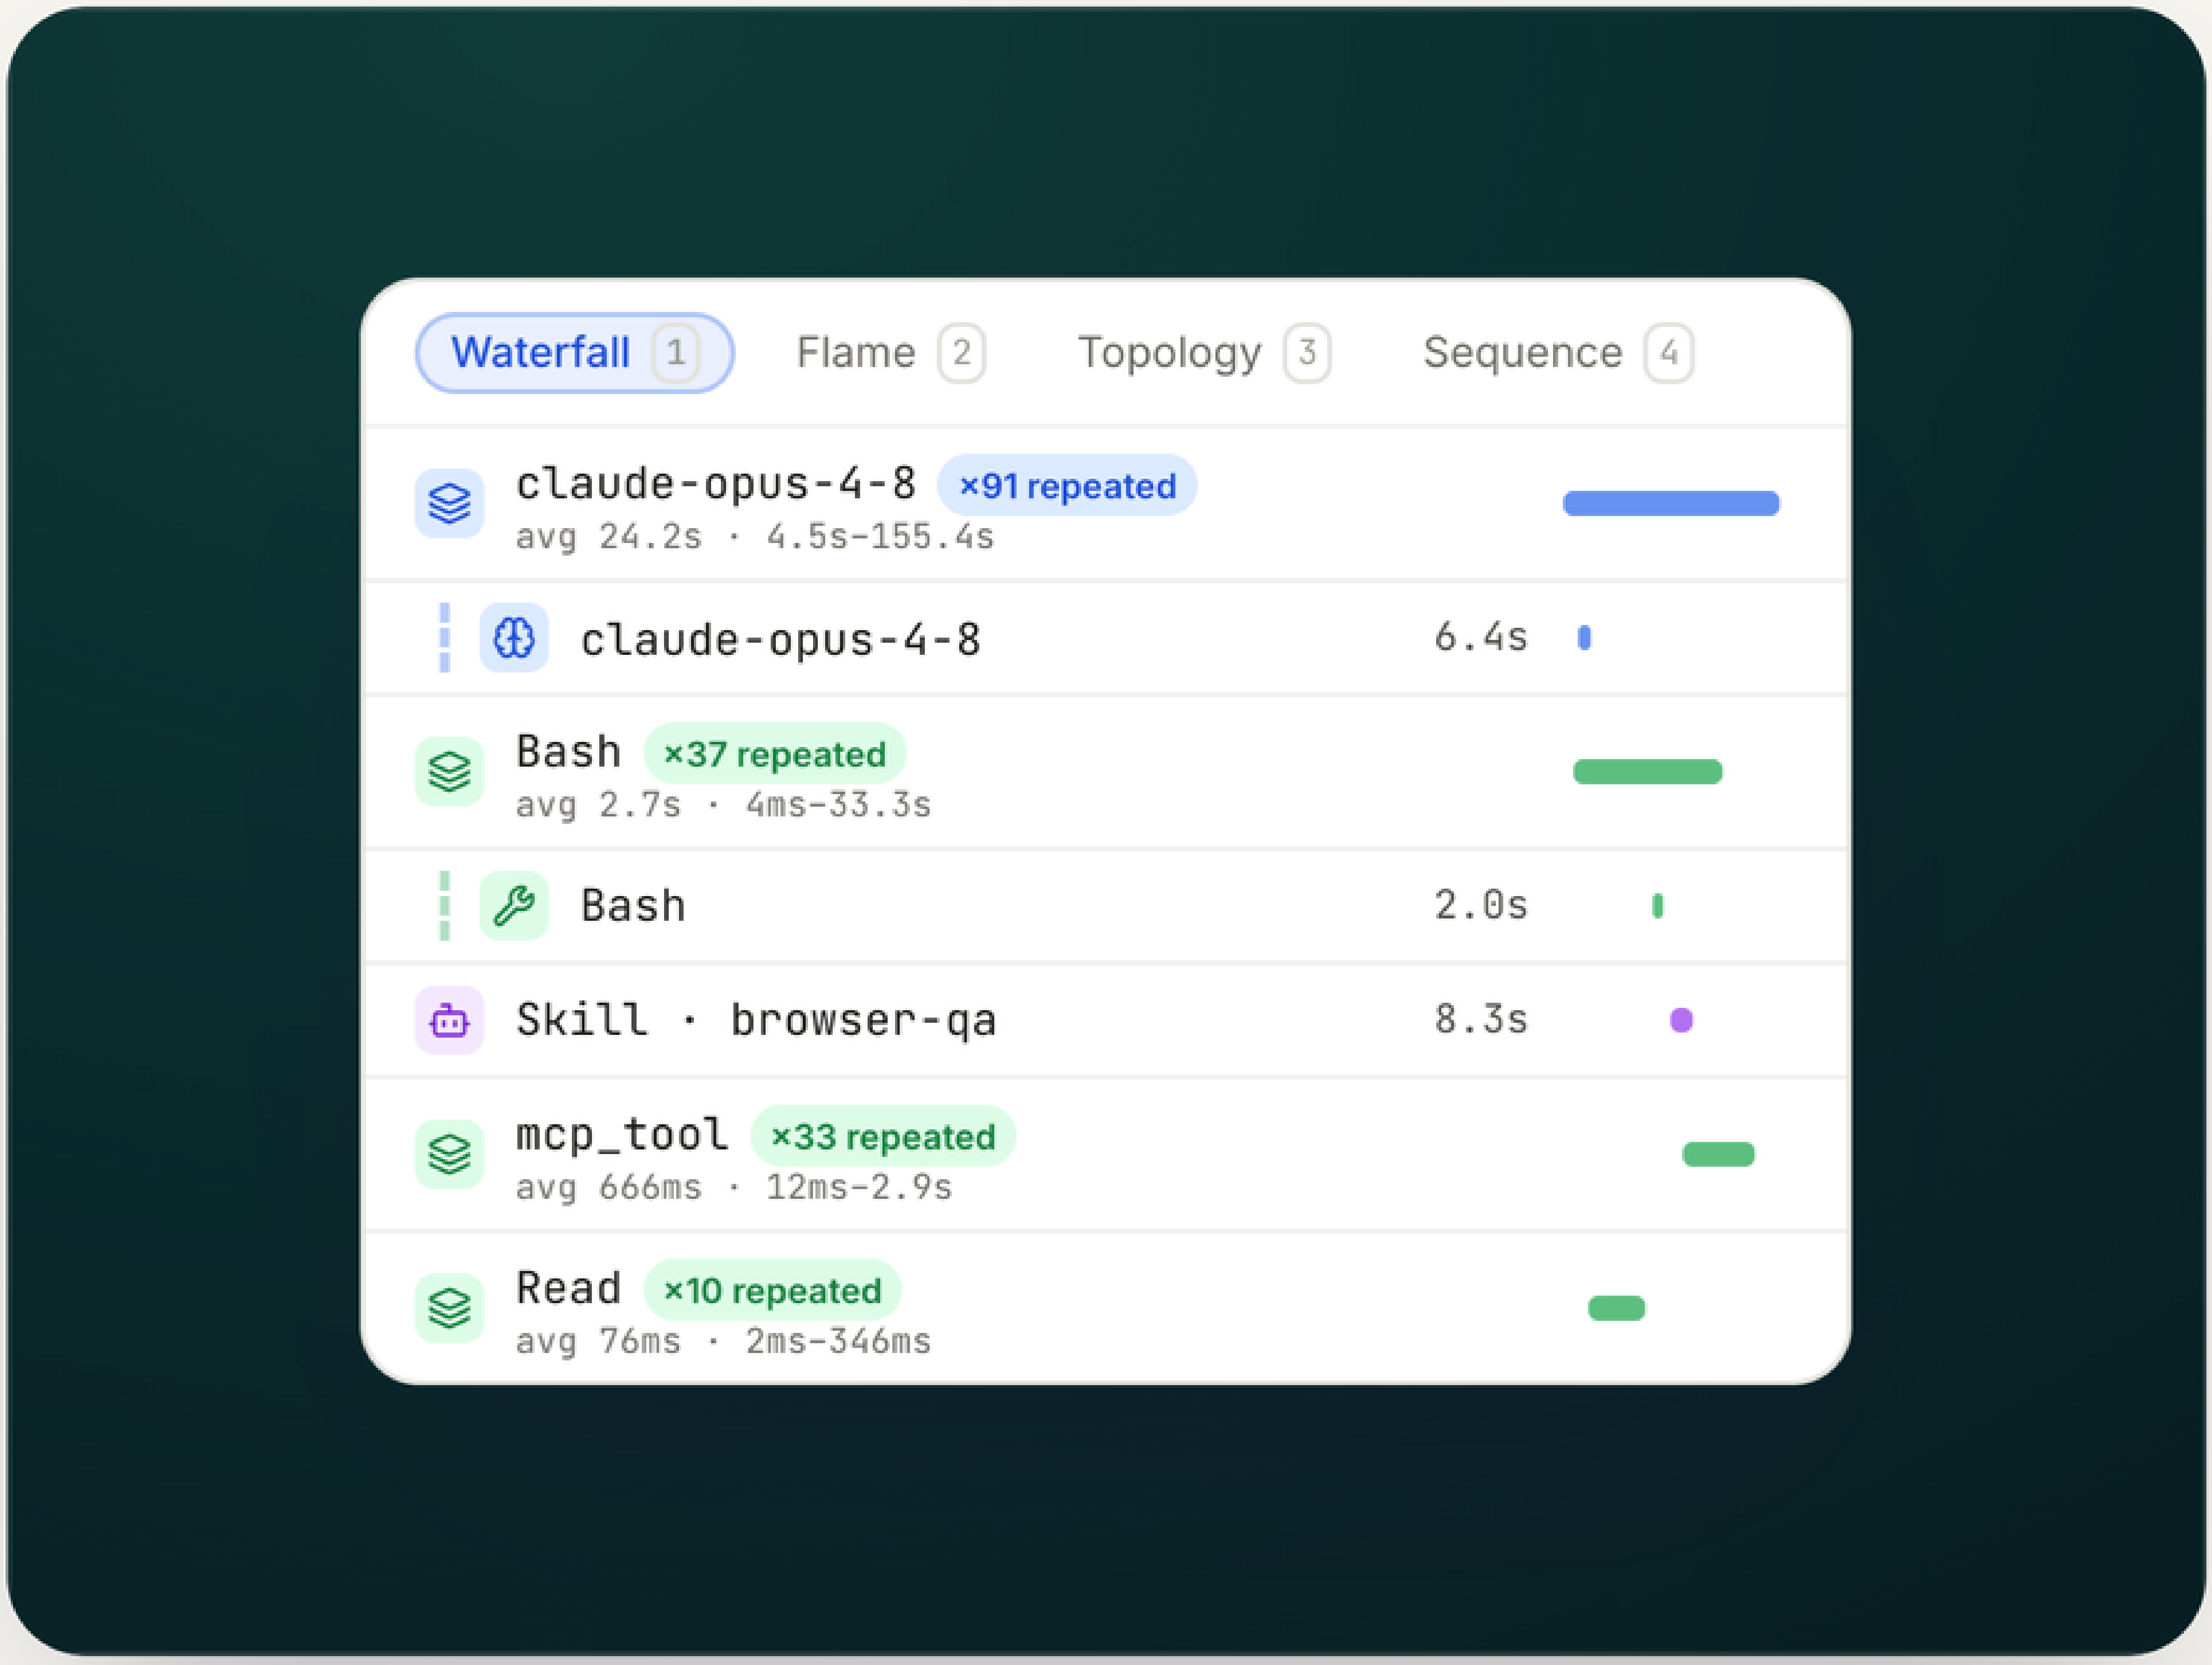

Tools, skills, MCP

Every Bash command, file edit, skill invocation, and MCP call is a span with a duration, and full input and output details.

```
claude code                                    mcp: langwatch

> where did my tokens go last week? check
langwatch

● langwatch · query_analytics(last 7d) ✓

  41 sessions, $212 theoretical. 62% is one
  pattern: the test suite re-runs in full
  instead of the failing shard.

> fix .claude/commands/test.md to target the
shard
```

Claude Code can query it back

The MCP server exposes your traces to the agent itself: it inspects last week’s runs, finds where tokens go, and edits its own workflow.

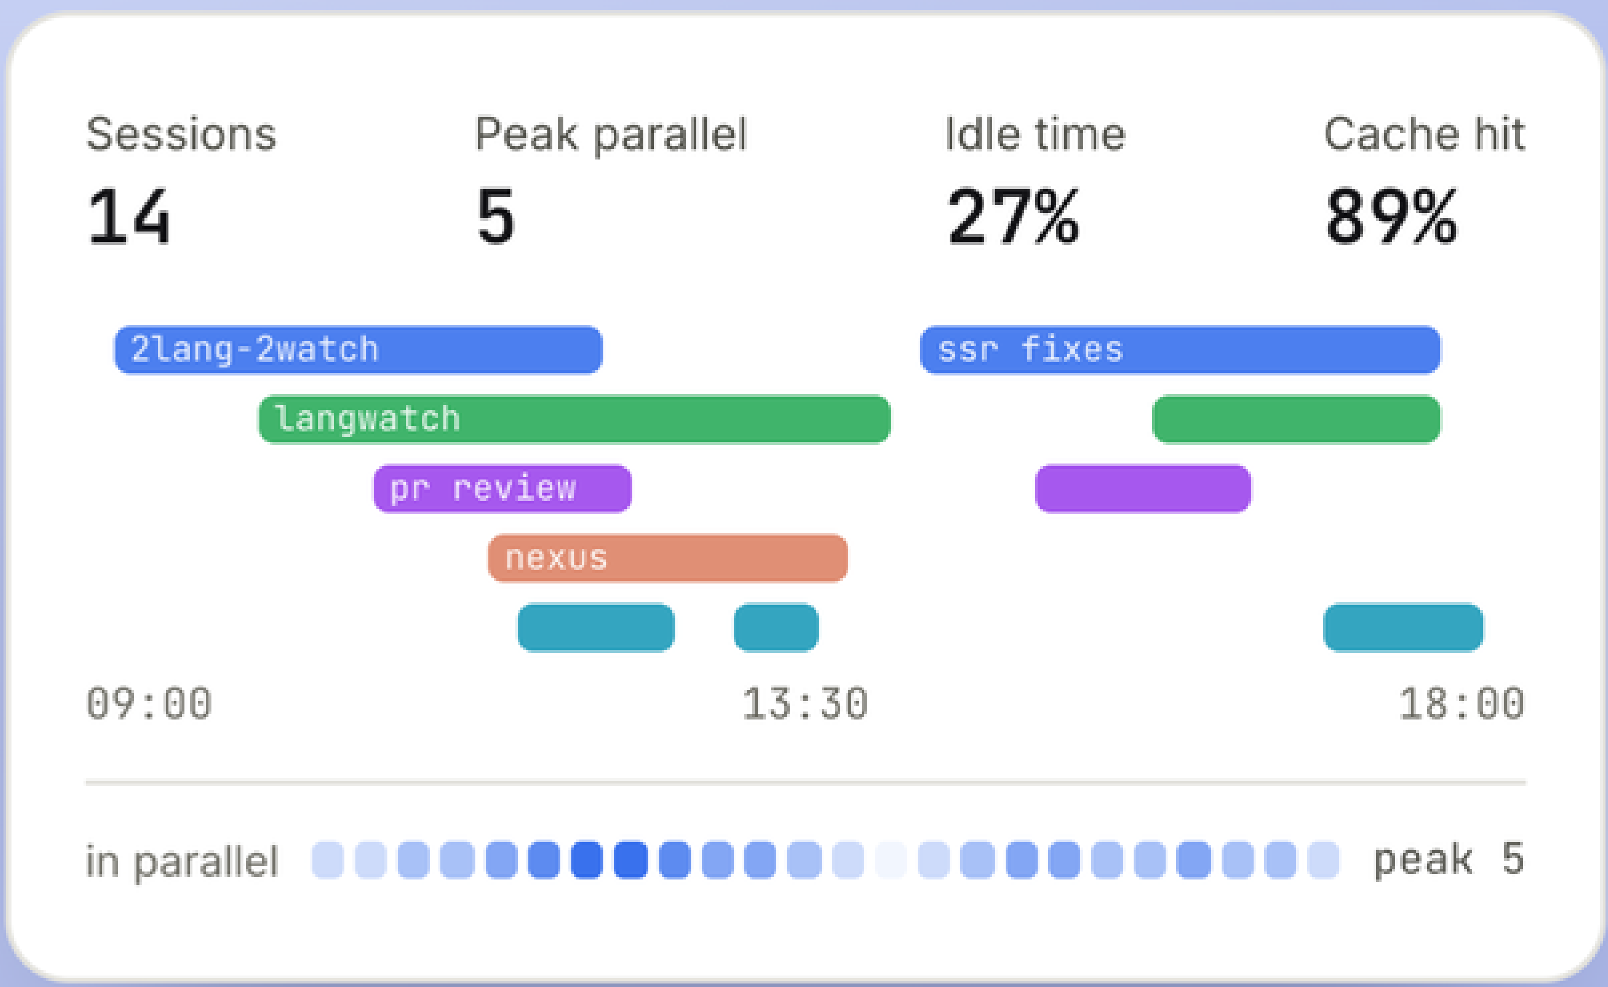

Plot any analytics

Any metric on the trace can go on a chart: sessions running in parallel, cost per session, idle time while it waits on your reply.

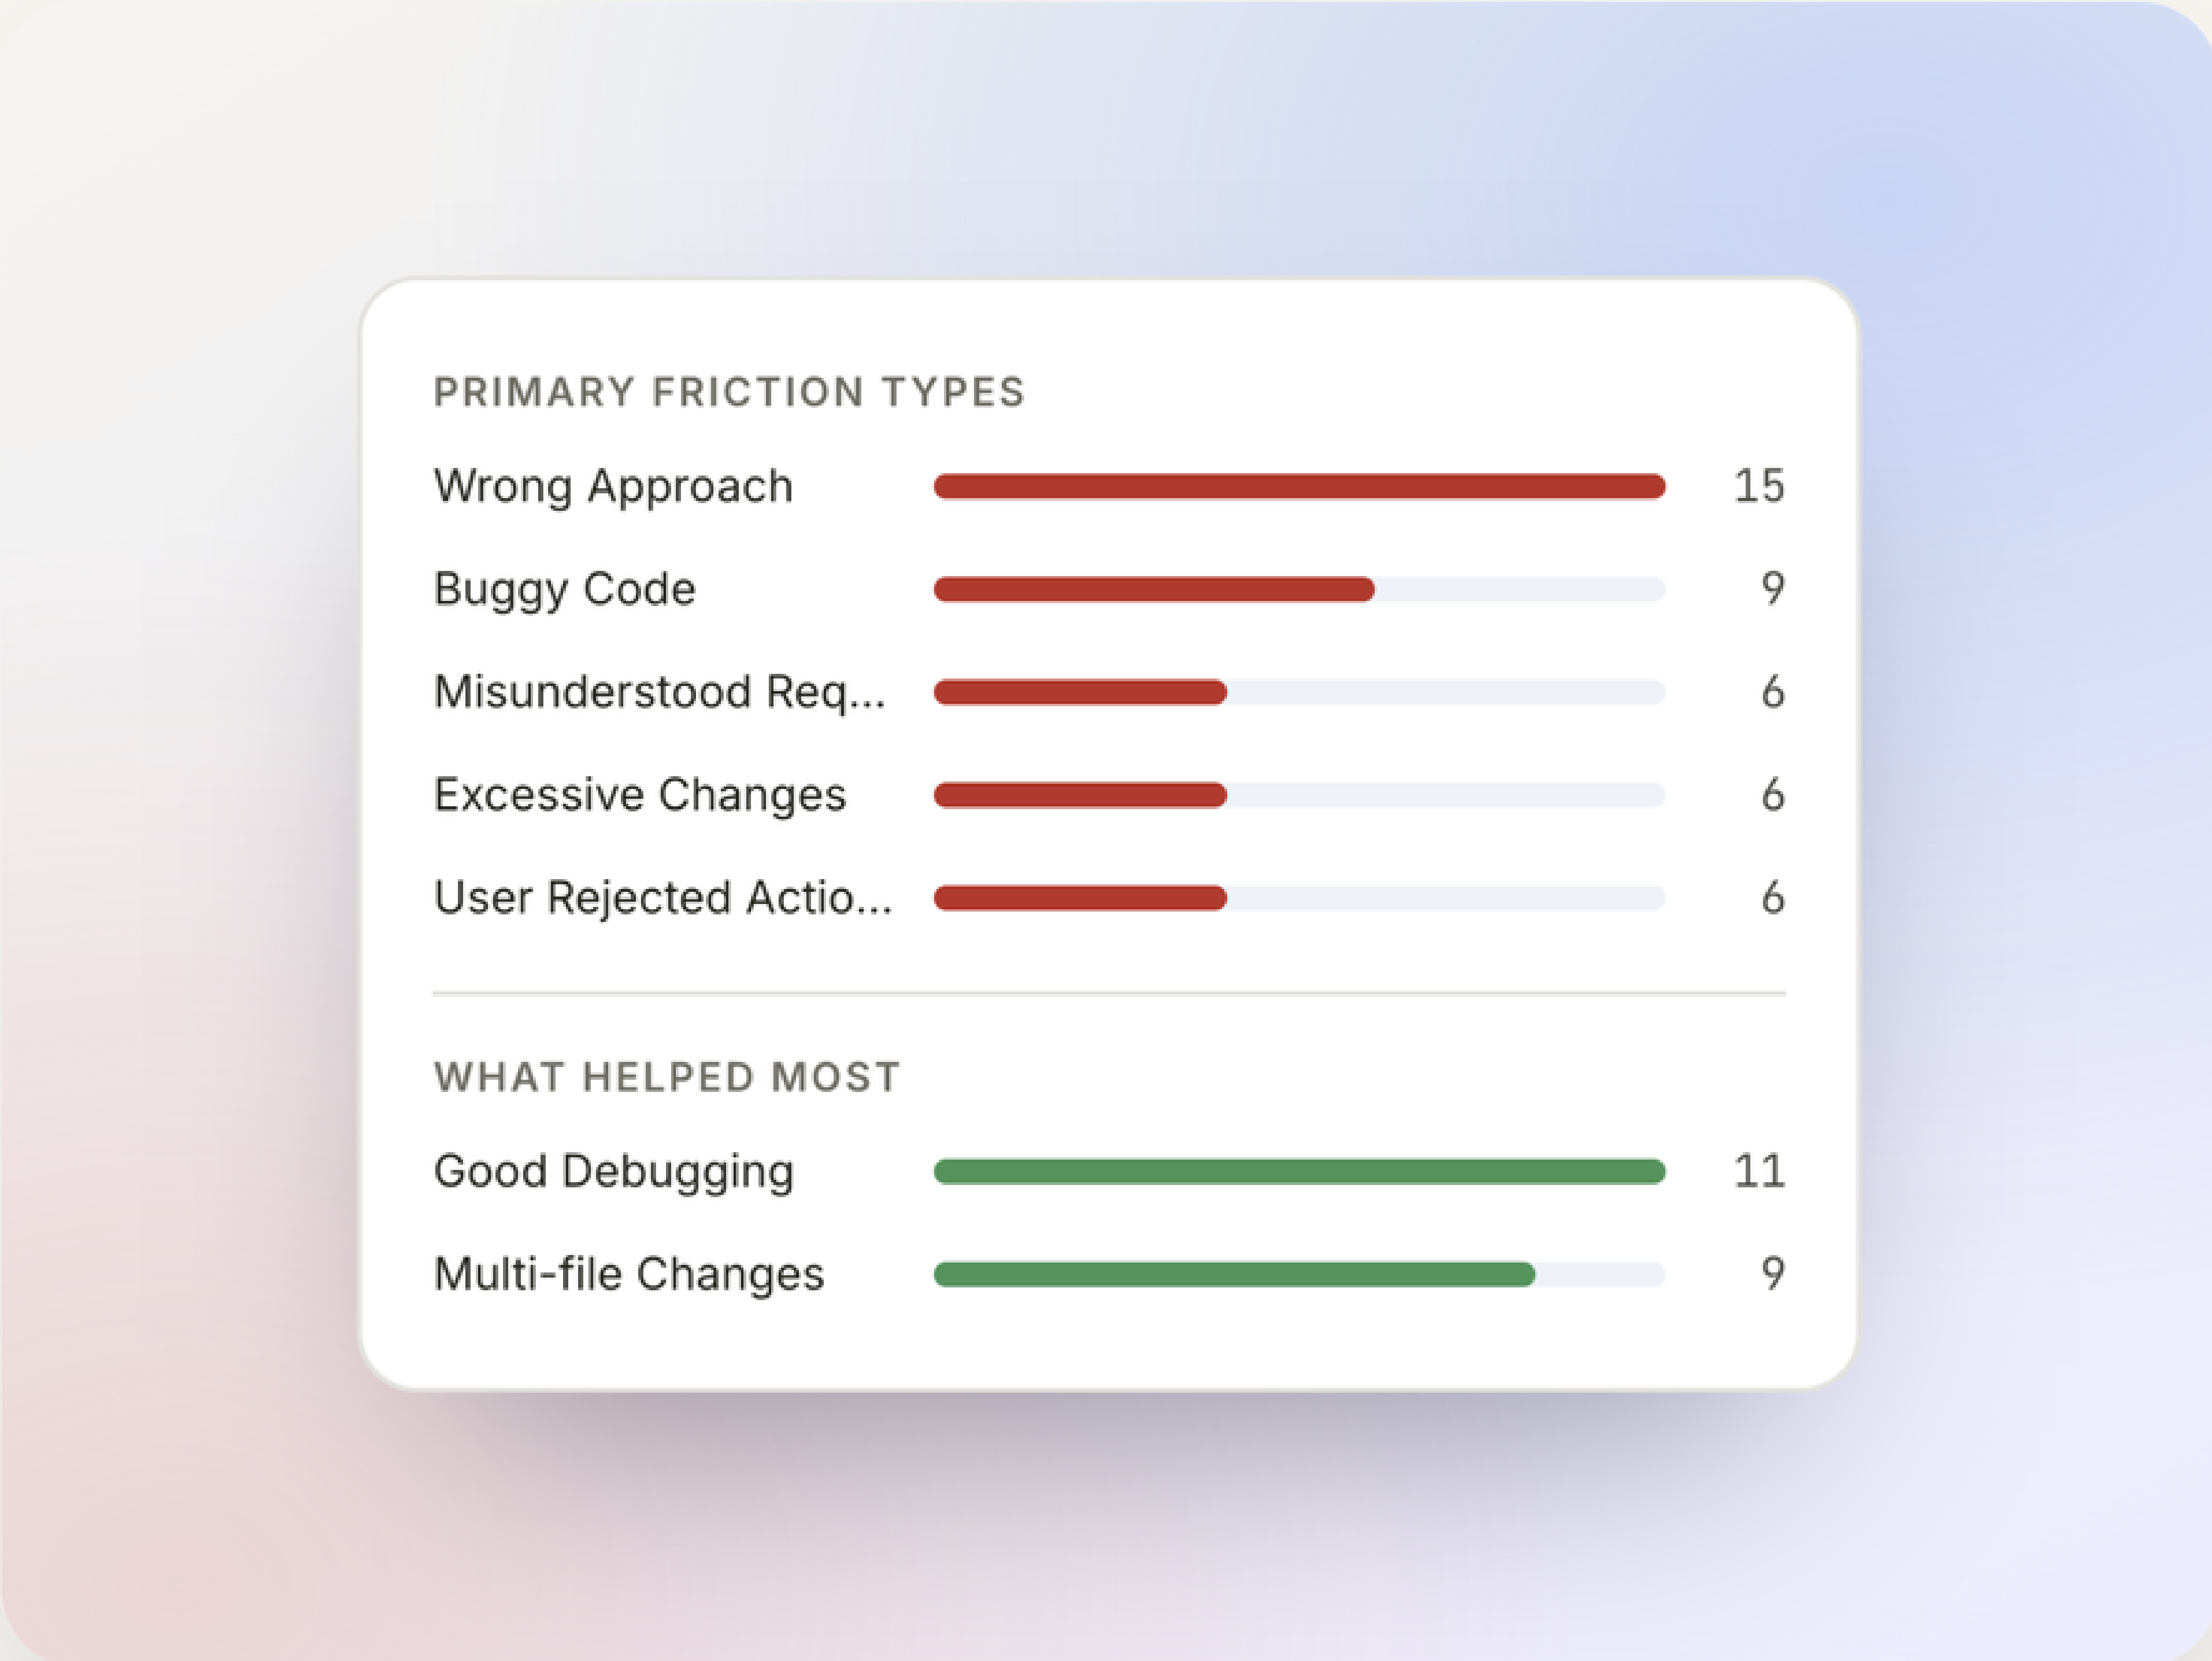

Insights to improve your skills

Run Claude over your own trace history: what you asked for, what helped most, and the friction patterns where sessions go wrong.

```
$ npx langwatch claude
```

Free for individual use.

langwatch.ai/claude-code-usage

LangWatch