



Track your Claude Code usage



Lang**Watch**




Tokens 124 in · 7.9K out

Models **A** claude-fable-5 +1

Input	124
Output	7,852
Cache read	8,398,963
Cache write	8,044
ephemeral 5m	0
ephemeral 1h	8,044
Total	8,414,983

Cache hit tracking

Cache reads and writes are separate token classes on every span, down to the 5m and 1h TTL split on writes.



Theoretical vs billed costs

Exact cost numbers, straight from the API response. See how much you would be spending if you weren't on an individual or team plan.



Tools, skills, MCP

Every Bash command, file edit, skill invocation, and MCP call is a span with a duration. ×37 repeated unfolds into each one.

claude code

mcp: langwatch

> where did my tokens go last week? check langwatch


• langwatch · query_analytics(last 7d) ✓

41 sessions, \$212 theoretical. 62% is one pattern: the test suite re-runs in full instead of the failing shard.

> fix .claude/commands/test.md to target the shard

Claude Code can query it back

The MCP server exposes your traces to the agent itself: it inspects last week's runs, finds where tokens go, and edits its own workflow.



Plot any analytics

Any metric on the trace can go on a chart:
sessions running in parallel, cost per session, idle
time while it waits on your reply.

PRIMARY FRICTION TYPES



WHAT HELPED MOST



Insights to improve your skills

Run Claude over your own trace history: what you asked for, what helped most, and the friction patterns where sessions go wrong.

```
$ npx langwatch claude
```

Free for individual use.

langwatch.ai/claude-code-usage